

UCR

CS161 – Design and Architecture of Computer Systems

Week 5 - Discussion

*some slides adapted from:
[Prof Daniel Wong UCR – EE/CS](#)

UNIVERSITY OF CALIFORNIA, RIVERSIDE

Memory vs Storage

- > Disk (aka: **Secondary Storage**):
 - > used to store data of **all your programs** and everything else on the computer
 - > often referred to as (when buying a computer):
 - > “Storage” > “Solid State Drive (SSD)”
 - > “Hard Drive” > “Flash memory” (for phones)
- > Main memory:
 - > used to hold **programs while they are running**
 - > can be thought of as a “cache” for disk storage
 - > often referred to as (when buying a computer):
 - > “Memory” > “RAM”

Memory vs Storage

➤ Main memory:

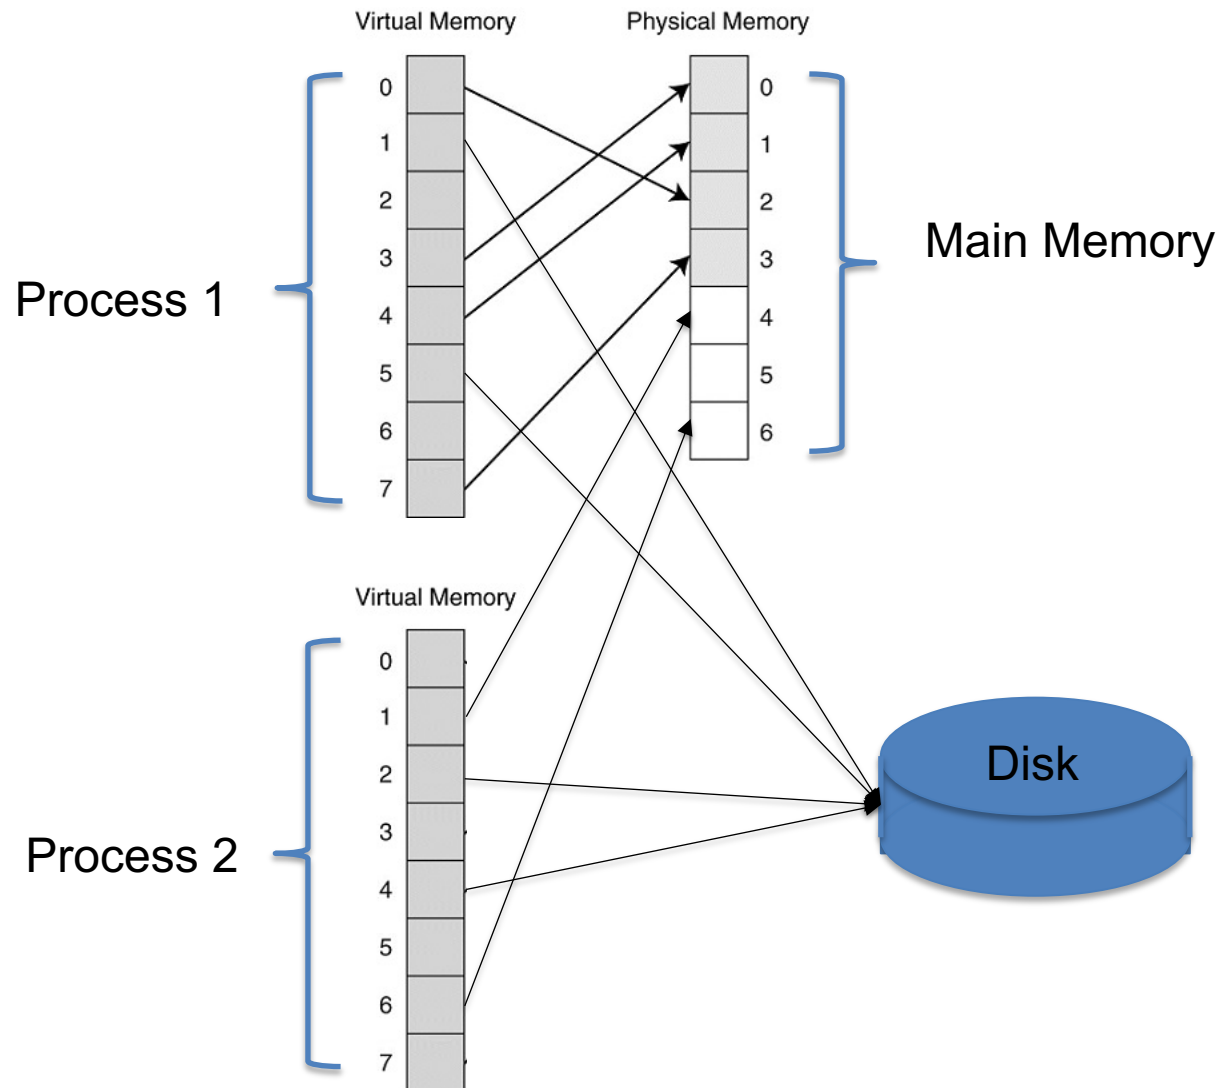
- used to hold programs while they are running
- so programs are limited by size of memory??
- and how do we switch programs running in memory??



Virtual Memory: Big Picture

- Each program thinks it has entire computer disk to run on (“Virtual Memory”)
- However, each program runs in shared main memory (“Physical Memory”)
- But: main memory \ll disk
- And: multiple programs running at once (sharing main memory)
- So we have to convert virtual memory to actual physical memory or actual disk
- Thus: **main memory is a “cache” for disk**

Shared Physical (Main) Memory




Pages

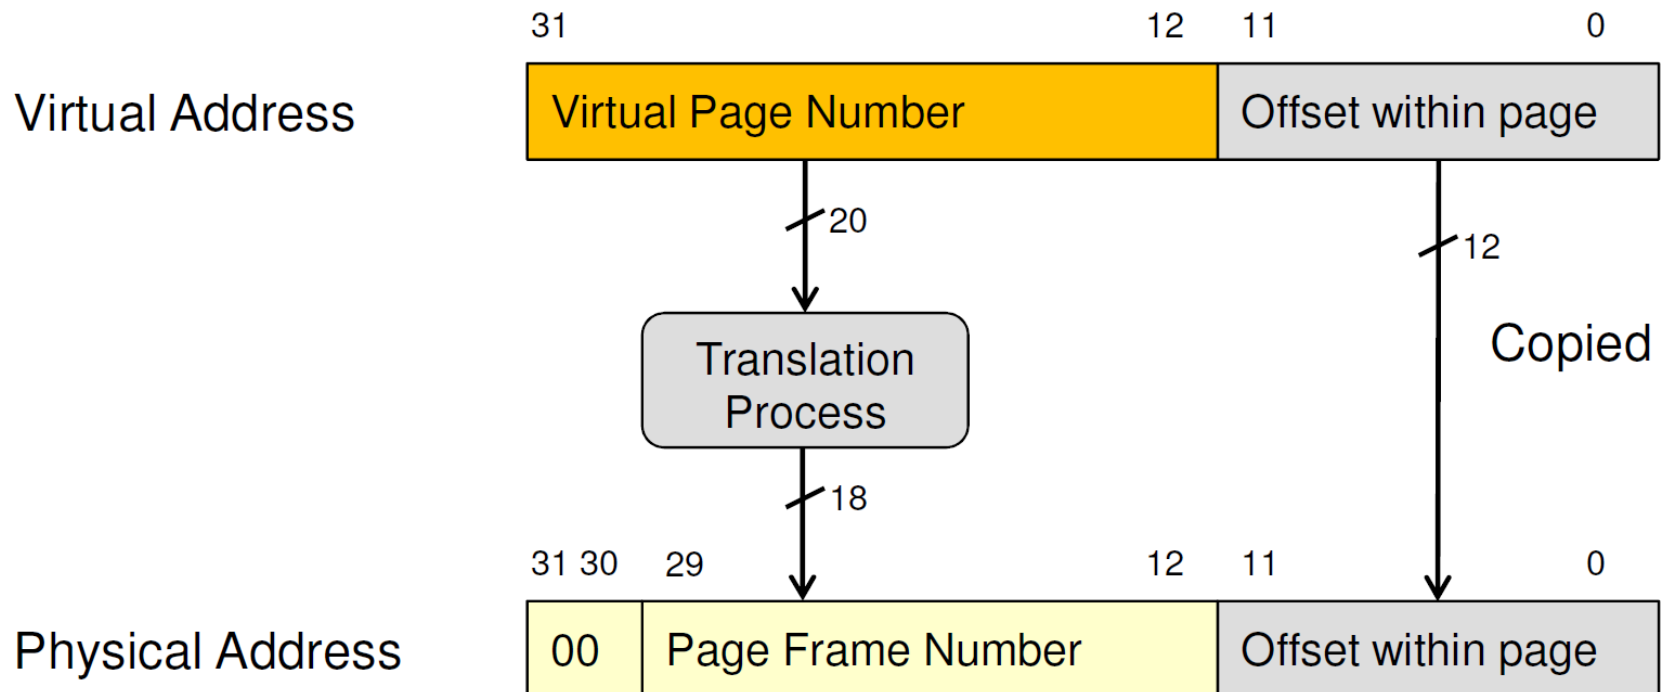
- Divide memory into equal sized “chunks” or pages (i.e., 4KB each)
- Any chunk of Virtual Memory can be assigned to any chunk of Physical Memory
- A page
 - a fixed-length contiguous block of virtual memory, described by a single entry in the page table
 - It is the smallest unit of data for memory management in a virtual memory operating system.
- page frame
 - the smallest fixed-length contiguous block of physical memory into which memory pages are mapped by the operating system.
- A transfer of pages between main memory and an auxiliary store, such as a hard disk drive, is referred to as paging or swapping

Page Fault

- Page missing in Main Memory
- Pages not in Main Memory are on disk:
 - swap-in/out




Mapping Virtual to Physical Address



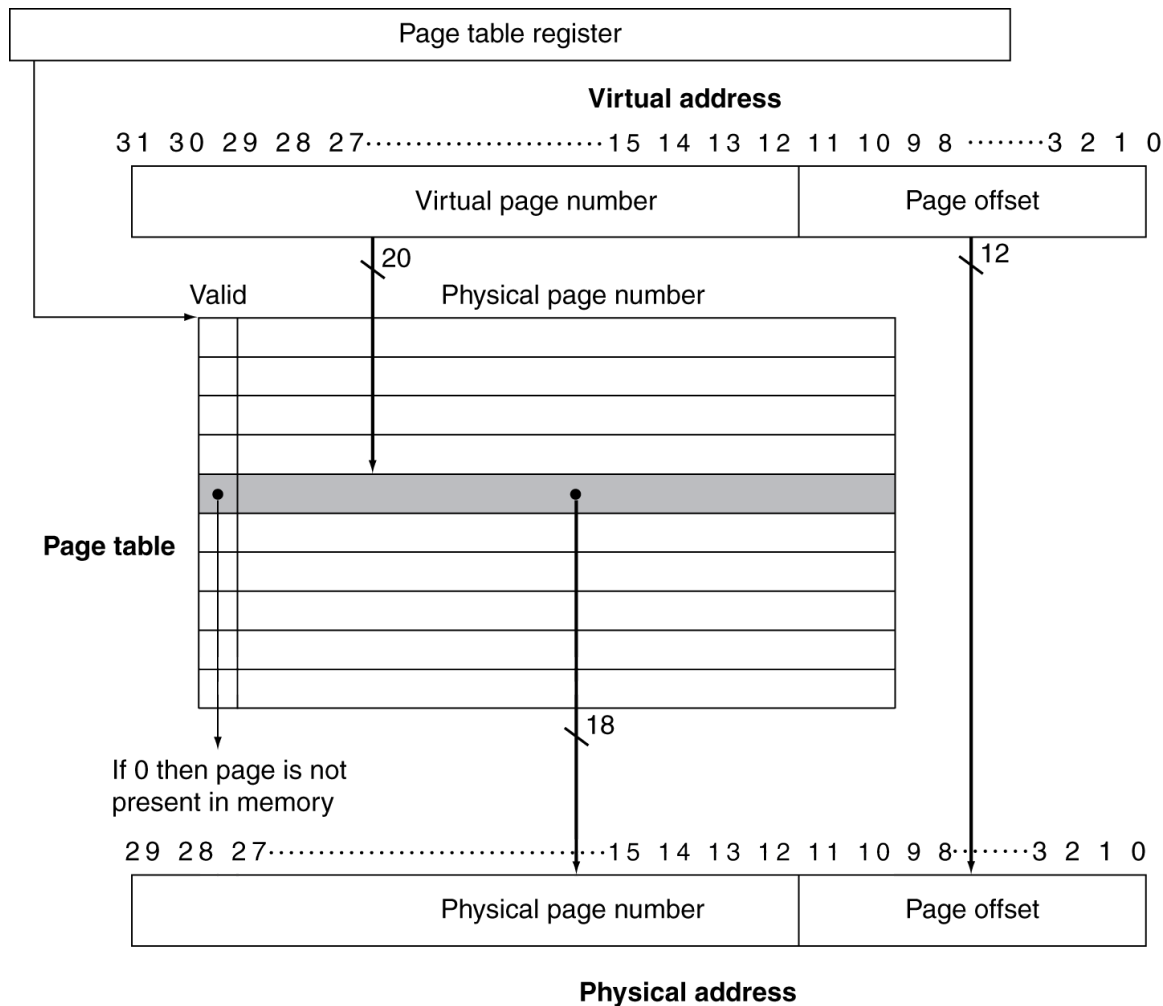
Address Translation: Page Table

- › A page table is a
 - › data structure containing mapping of virtual to physical pages
 - › each process has its own page table

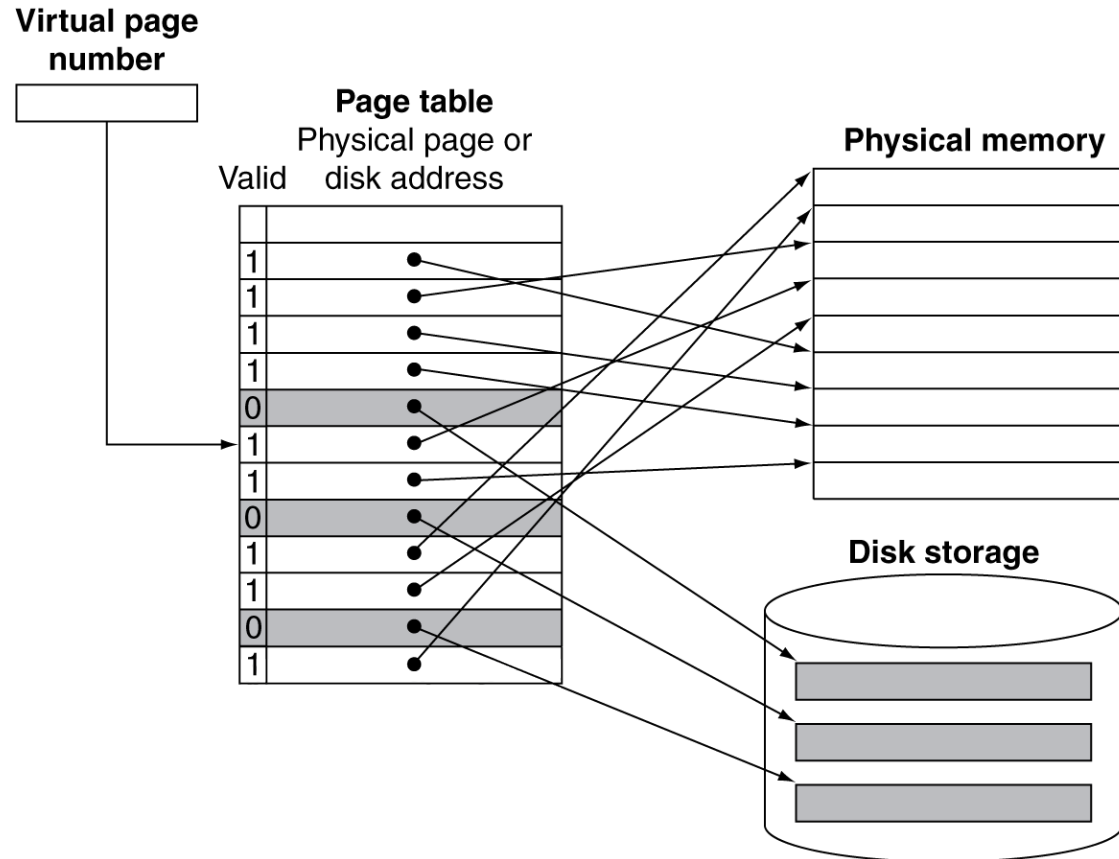


Page Table

- Page table translates address



Mapping Pages to Storage



Optimizing VM

› Page Table too big!

- › virtual address space = 4GB = $2^2 2^{30} \text{ B} = 2^{32} \text{ B}$
- › page size = 4 KB = $2^2 2^{10} \text{ B} = 2^{12} \text{ B}$
- › PTE size = 4 B
 - › number page table entries = $2^{32} - 2^{12} = 2^{20}$
 - › **Page table size = 4 B * $2^{20} = 4\text{MB}$**
(just for Page Table of single process)

› Virtual Memory too slow!


- › Requires two memory accesses.
 - › One to access page table to get the memory address
 - › Another to get the real data

Multi-level Page Table

- Problem: 1-level page table too expensive for large virtual address space
 - Need to reduce the size of the page table
- Solution: Multi-level page table, Paging page tables, etc.
 - To create small page tables for virtual memory
 - The virtual address is now split into multiple chunks to index a page table “tree”

Multi-level Page Table


- Virtual page number broken into fields to index each level of multi-level page table



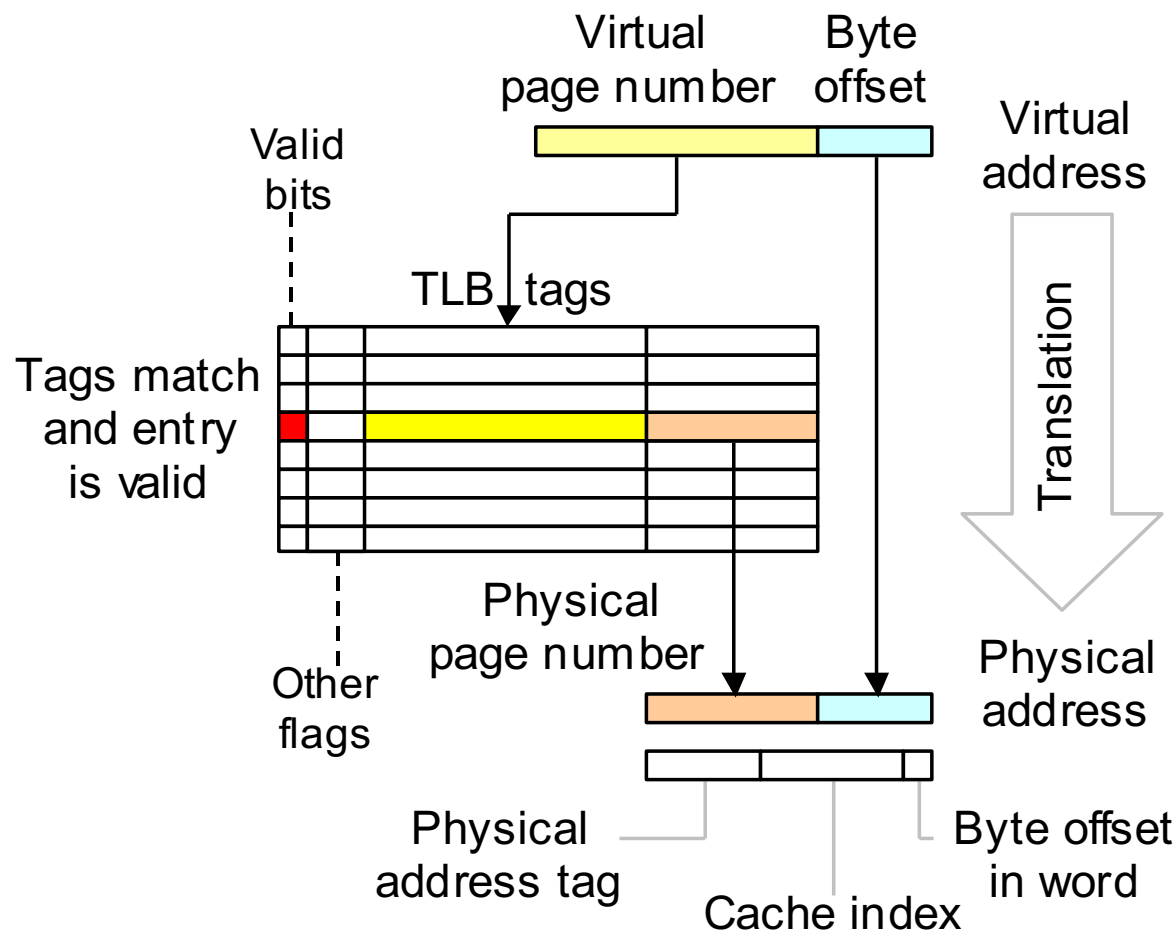
Fast Address Translation

- Problem: Virtual Memory requires two memory accesses!
 - one to translate Virtual Address into Physical Address (page table lookup)
 - one to transfer the actual data (cache hit)
 - But Page Table is in physical memory! => 2 main memory accesses!
- Solution: Why not create a cache of virtual to physical address translations to make translation fast? (smaller is faster)
- For historical reasons, such a “page table cache” is called a Translation Lookaside Buffer, or TLB

Fast Translation Using a TLB




TLB Translation



Virtual-to-physical address translation by a TLB and how the resulting physical address is used to access the cache memory.


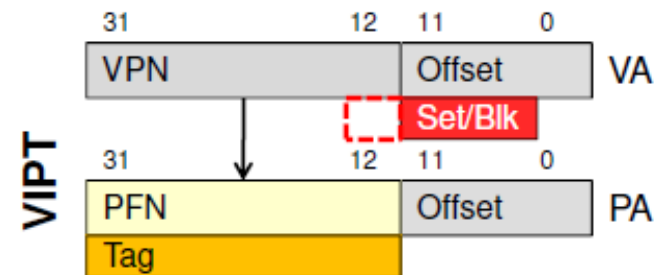

TLB and Cache Interaction



- If cache tag uses physical address
 - Need to translate before cache lookup
- Physically Indexed, Physically Tagged

Cache Index/Tag Options

- Physically indexed, physically tagged (PIPT)
 - Wait for full address translation
 - Then use physical address for both indexing and tag comparison
- Virtually indexed, physically tagged (VIPT)
 - Use portion of the virtual address for indexing then wait for address translation and use physical address for tag comparisons
 - Easiest when index portion of virtual address w/in offset (page size) address bits, otherwise aliasing may occur
- Virtually indexed, virtually tagged (VIVT)
 - Use virtual address for both indexing and tagging...No TLB access unless cache miss
 - Requires invalidation of cache lines on context switch or use of process ID as part of tags



Summary

- › Virtual Memory overcomes main memory size limitations
- › VM supported through Page Tables
- › Multi-level Page Tables enables smaller page tables in memory
- › TLB enables fast address translation

Example 1

1. The following list provides parameters of a virtual memory system:

Virtual Address (bits)	Physical DRAM Installed	Page Size	PTE Size (bytes)
43	16 GB	4 KB	4

a) For a single-level page table, how many page table entries (PTEs) are needed?

Example 1

1. The following list provides parameters of a virtual memory system:

Virtual Address (bits)	Physical DRAM Installed	Page Size	PTE Size (bytes)
43	16 GB	4 KB	4

a) For a single-level page table, how many page table entries (PTEs) are needed?

virtual address = 43 bits

page size = 4 KB = $2^2 2^{10} \text{B} = 2^{12} \text{B} \rightarrow$ offset bits = 12

virtual page number (VPN) bits = virtual address bits – offset bits
= 43 – 12
= 31

entries in page table = $2^{31} = 2^1 2^{30} = 2 \text{ G entries}$

Example 1

1. The following list provides parameters of a virtual memory system:

Virtual Address (bits)	Physical DRAM Installed	Page Size	PTE Size (bytes)
43	16 GB	4 KB	4

b) How much physical memory is needed for storing the page table?

Example 1

1. The following list provides parameters of a virtual memory system:

Virtual Address (bits)	Physical DRAM Installed	Page Size	PTE Size (bytes)
43	16 GB	4 KB	4

b) How much physical memory is needed for storing the page table?

$$\begin{aligned}\text{page table size} &= (\text{PTE size in bytes}) * (\text{number entries in page table}) \\ &= (2 \text{ G}) * (4 \text{ B}) \\ &= (2^{31}) * (2^2 \text{ B}) \\ &= 2^{33} \text{ B} \\ &= 2^3 2^{30} \text{ B} \\ &= 8 \text{ GB}\end{aligned}$$

Example 1

1. The following list provides parameters of a virtual memory system:

Virtual Address (bits)	Physical DRAM Installed	Page Size	PTE Size (bytes)
43	16 GB	4 KB	4

c) What is the size of secondary storage?

Example 1

1. The following list provides parameters of a virtual memory system:

Virtual Address (bits)	Physical DRAM Installed	Page Size	PTE Size (bytes)
43	16 GB	4 KB	4

c) What is the size of secondary storage?

$$\begin{aligned}
 \text{secondary storage size (in bytes)} &= 2^{\text{virtual address bits}} \text{ bytes} \\
 &= 2^{43} \text{ bytes} \\
 &= 2^3 2^{40} \text{ bytes} \\
 &= 8 \text{ TB}
 \end{aligned}$$

Example 1

1. The following list provides parameters of a virtual memory system:

Virtual Address (bits)	Physical DRAM Installed	Page Size	PTE Size (bytes)
43	16 GB	4 KB	4

d) How many bits of physical address are needed to address DRAM (main memory)?

Example 1

1. The following list provides parameters of a virtual memory system:

Virtual Address (bits)	Physical DRAM Installed	Page Size	PTE Size (bytes)
43	16 GB	4 KB	4

d) How many bits of physical address are needed to address DRAM (main memory)?

$$\begin{aligned}\text{physical DRAM size} &= 2^{\text{physical address bits}} \text{ bytes} \\ &= 16 \text{ GB} \\ &= 2^4 2^{30} \text{ B} \\ &= 2^{34} \text{ B.} \rightarrow 34 \text{ bits needed}\end{aligned}$$

Example 2

2. Using a multi-level page table with the following list provides parameters of a virtual memory system:

Virtual Address (bits)	Physical DRAM Installed	Page Size	PTE Size (bytes)
43	16 GB	4 KB	4

a) How many levels of page tables will be needed in this case?

Example 2

2. Using a multi-level page table with the following list provides parameters of a virtual memory system:

Virtual Address (bits)	Physical DRAM Installed	Page Size	PTE Size (bytes)
43	16 GB	4 KB	4

a) How many levels of page tables will be needed in this case?

$$\begin{aligned} \text{number of page table entries that fit in a page} &= \frac{\text{page size}}{\text{PTE size}} = \frac{4 \text{ KB}}{4 \text{ B}} = 1 \text{ K} \\ &= 2^{10} \text{ entries in a page} \end{aligned}$$

→ so we need 10 bits to find the PTE in a page

$$\text{number of levels} = \frac{\text{bits of single level page table entries}}{\text{bits of multilevel page table entries}} = \frac{31}{10} = 3.1 = 4 \text{ levels}$$

Example 2

2. Using a multi-level page table with the following list provides parameters of a virtual memory system:

Virtual Address (bits)	Physical DRAM Installed	Page Size	PTE Size (bytes)
43	16 GB	4 KB	4

b) How many memory references are needed for address translation if missing in TLB?

Example 2

2. Using a multi-level page table with the following list provides parameters of a virtual memory system:

Virtual Address (bits)	Physical DRAM Installed	Page Size	PTE Size (bytes)
43	16 GB	4 KB	4

b) How many memory references are needed for address translation if missing in TLB?

if there is a miss in the TLB for address translation, then 4 memory references are required, one for each level

Example 3

3. Given the following parity bits, data bits, and field coverage in a Hamming ECC code for eight data bits (i.e., a [12,8] encoding):

Bit position		1	2	3	4	5	6	7	8	9	10	11	12
Encoded data bits		p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8
Parity bit coverage	p1	X		X		X		X		X		X	
	p2		X	X			X	X			X	X	
	p4				X	X	X	X					X
	p8								X	X	X	X	X

a) decode the bits for the value 0x9A.

Example 3

3. Given the following parity bits, data bits, and field coverage in a Hamming ECC code for eight data bits (i.e., a [12,8] encoding):

Bit position		1	2	3	4	5	6	7	8	9	10	11	12
Encoded data bits		p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8
Parity bit coverage	p1	X		X		X		X		X		X	
	p2		X	X			X	X			X	X	
	p4				X	X	X	X					X
	p8								X	X	X	X	X

a) decode the bits for the value 0x9A.

$$0x9A = 1001\ 1010 = 10011010$$

So there's eight data bits (8) and we're using 12 bit parity encoding, so we will have 4 bits for parity bits, which we will put in spots 1, 2, 4, and 8, so we leave a blank there for now:

__1__001__1010

p1 checks every other spot (1 position in between), so positions 1,3,5,7,9, and 11, the positions are highlighted in red:

__1__001__1010

since there are four 1's highlighted, then we should set the first bit (i.e., the **p1** bit) to 0 to keep the group's parity at an even number:

0__1__001__1010

Example 3

3. Given the following parity bits, data bits, and field coverage in a Hamming ECC code for eight data bits (i.e., a [12,8] encoding):

Bit position	1	2	3	4	5	6	7	8	9	10	11	12
Encoded data bits	p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8
Parity bit coverage	p1	X		X		X		X		X		
	p2		X	X			X	X			X	X
	p4				X	X	X	X				X
	p8							X	X	X	X	X

a) decode the bits for the value 0x9A.

next we look at the case for **p2**, checking every two bits (highlighted in red), so we're checking positions 2,3,6,7,10,and 11:

0_1_001_1010

since there are three 1's highlighted, then we should set the second bit (i.e., the **p2** bit) to 1 to make the group's parity at an even number:

011_001_1010

next we look at the case for **p4**, checking every four bits (highlighted in red), so we're checking positions 4,5,6,7,and 12:

011_001_1010

Example 3

3. Given the following parity bits, data bits, and field coverage in a Hamming ECC code for eight data bits (i.e., a [12,8] encoding):

Bit position	1	2	3	4	5	6	7	8	9	10	11	12
Encoded data bits	p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8
Parity bit coverage	p1	X		X		X		X		X		
	p2		X	X		X	X			X	X	
	p4				X	X	X	X				X
	p8							X	X	X	X	X

a) decode the bits for the value 0X9A.

since there is only one 1 highlighted, then we should set the fourth bit (i.e., the **p4** bit) to 1 to make the group's parity at an even number:

0111001_1010

finally, we look at the case for **p8**, checking every eight bits (highlighted in red), so we're checking positions 8,9,10,11,and 12:

0111001_1010

since there are two 1's highlighted, then we should set the eighth bit (i.e., the **p8** bit) to 0 to keep the group's parity at an even number:

011100101010

this leaves us with the following final code word:

011100101010

Example 3

3. Given the following parity bits, data bits, and field coverage in a Hamming ECC code for eight data bits (i.e., a [12,8] encoding):

Bit position	1	2	3	4	5	6	7	8	9	10	11	12
Encoded data bits	p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8
Parity bit coverage	p1	X		X		X		X		X		
	p2		X	X			X	X		X	X	
	p4				X	X	X	X				X
	p8							X	X	X	X	X

a) decode the bits for the value 0x9A.

but let's assume that the following was received instead (bit 10 was flipped):

011100101**1**10

so to check if there was an error, we would check the parity and compare them with their respective parity bits:

p1: 0**1**1100**1**0**1**110 → four 1's (even) without parity bit and parity bit is 0 (even) → no error

p2: 0**1**1100**1**0**1**110 → four 1's (even) without parity bit and parity bit is 1 (odd) → **error**

p4: 011**1**00**1**01110 → one 1's (odd) without parity bit and parity bit is 1 (odd) → no error

p8: 0111001**0**1110 → three 1's (odd) without parity bit and parity bit is 0 (even) → **error**

since parity bits 2 (p2) and 8 (p8) are incorrect, then bit 10 must be wrong (2+8=10)

011100101**1**10 → 011100101**0**10